

# Introduction to Reinforcement Learning

Mattia Pellegrino, Ph.D Fellow  
[mattia.pellegrino@unipr.it](mailto:mattia.pellegrino@unipr.it)



# Summary

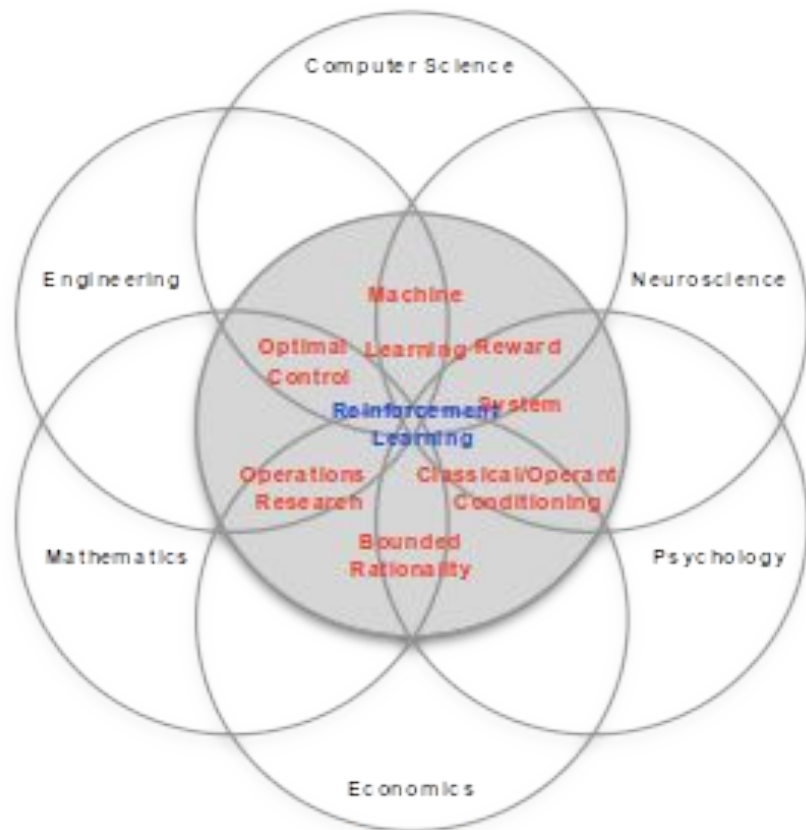
- Information about the course
- About Reinforcement Learning
- The Reinforcement Learning challenges
- RL Agent
- Reinforcement Learning Problems

# COURSE INFORMATION

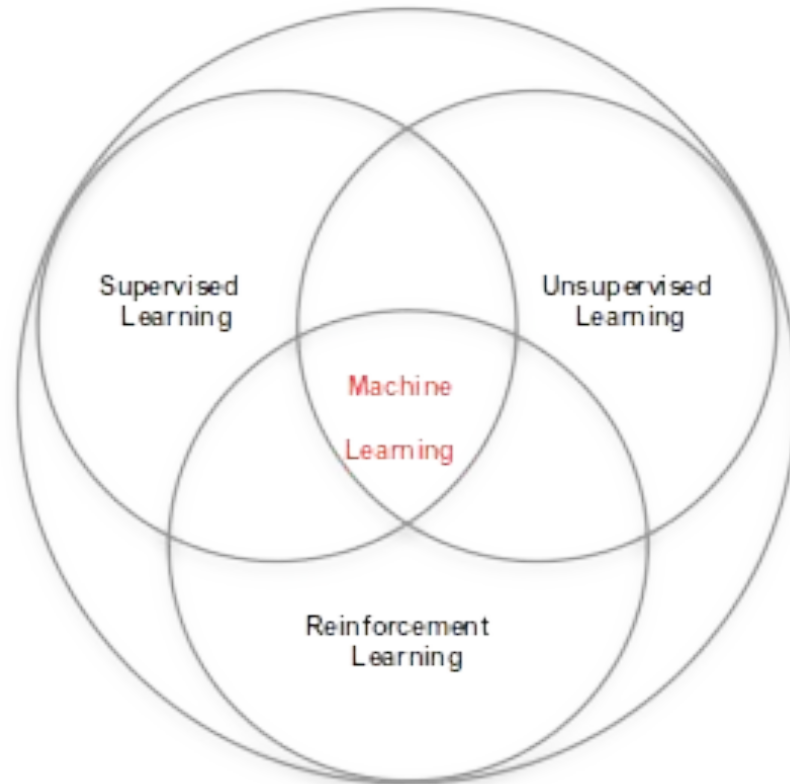
- Contact me: [mattia.pellegrino@unipr.it](mailto:mattia.pellegrino@unipr.it)
- **TextBooks:**
  - Montague, P. Read. "Reinforcement learning: an introduction, by Sutton, RS and Barto, AG." Trends in cognitive sciences 3.9 (1999): 360. <http://webdocs.cs.ualberta.ca/~sutton/book/the-book.html>
  - Appress, "**Deep Reinforcement Learning with python**, by Nimish Sanghi", (2021)

# RL PLACEMENT

- All these branches try to solve the same problem: “**Decision Making**”
- Find the best actions’ combination to get the best result




# MACHINE LEARNING BRANCHES



# RL: PROPERTIES

- What makes RL different from other learning methodologies?
  - There is no supervisor. Instead, there is a *reward signal*
    - *How do we know if what we are doing is right or not?*
  - Reward is not instantaneous
    - *We realize that we have done a bad choice afterwards*
  - Time matters
  - Agent's actions affect the entire environment

# RL: A ROBOT LEARNS TO WALK



## RL: TRACKMANIA





# RL: BOXING



# RL: ATARI GAME



# RL: REWARD

- A reward  $R_t$  is a scalar feedback signal
  - It is just a number
- Indicates how good is the action choose by the agent at **timestep  $t$**
- The agent's goal is to maximize the total reward

RL is based on the *reward hypothesis*

**“All goals can be described by the maximization of expected cumulative reward”**

# RL: EXAMPLE OF REWARDS


- **A robot learns to walk**
  - + reward for forward motion
  - - reward for falling over
- **Trackmania**
  - + reward for forward motion
  - - reward to falling over
- **Boxing**
  - + reward to stand stand in place
  - - reward to falling over
  - + reward to hit the enemy
  - - reward to get hit
- **Play Atari Games**
  - +/- reward for increasing/decreasing score

# RL: SEQUENTIAL DECISION MAKING


- **Goal:** select actions to maximize total future reward
- Actions may have long term consequences
- Reward may be **delayed**
- It may be better to **sacrifice immediate reward** to gain more **long-term reward**
- Examples:
  - A financial investment (may take months to mature)
  - Blocking opponent moves (might help winning chances many moves from now)

# AGENT AND ENVIRONMENT

- We will use the current formalism



# AGENT AND ENVIRONMENT



- At each step  $t$  the agent:
  - Execute an action  $A_t$
  - Obtain an observation  $O_t$
  - Obtain a reward  $R_t$
- The environment
  - Receives an action  $A_t$
  - Releases an observation  $O_{t+1}$
  - Releases a Reward  $R_{t+1}$

# HISTORY - STATE

*History is the sequence of observation, action, and rewards (usually huge)*

$$H_t = O_1, R_1, A_1, \dots, A_t, O_t, R_t$$

- An algorithm is a **mapping** between the **history** and what **happens next**:
  - The agent selects actions
  - The environment selects observations/rewards


*The state is the information used to determine what happens next*

- It's a **synthesis** of what happened and we base on this because the history usually is too big to compute
- Formally, state is a function of the history

$$S_t = f(H_t)$$




# ENVIRONMENT STATE



- The ***environment state***  $S_t^e$  is the environment's private representation
- This state is **not** usually **visible** to the agent
- Sometimes contains **irrelevant** information

# AGENT STATE



- The **agent state**  $S_t^a$  is the agent's internal representation
- Whatever information the agent can use to choose the next action
- It can be any function of history:

$$S_t = f(H_t)$$

# INFORMATION STATE

*An information state (a.k.a. Markov state) contains all useful information from the history*

- A state  $S_t$  is a Markov state if and only if


$$\mathbb{P}[S_{t+1} \mid S_t] = \mathbb{P}[S_{t+1} \mid S_1, \dots, S_t]$$

- The **future** is **independent** of the **past** given present

$$H_{1:t} \longrightarrow S_t \longrightarrow H_{t+1:\infty}$$

- Once the **state** is **known**, the **history** is **irrelevant**
- The environment state  $S_t^e$  is a Markov State
- The history  $H_t$  is a Markov state

# FULLY OBSERVABLE ENVIRONMENTS



- Full observability: agent directly observes environment state

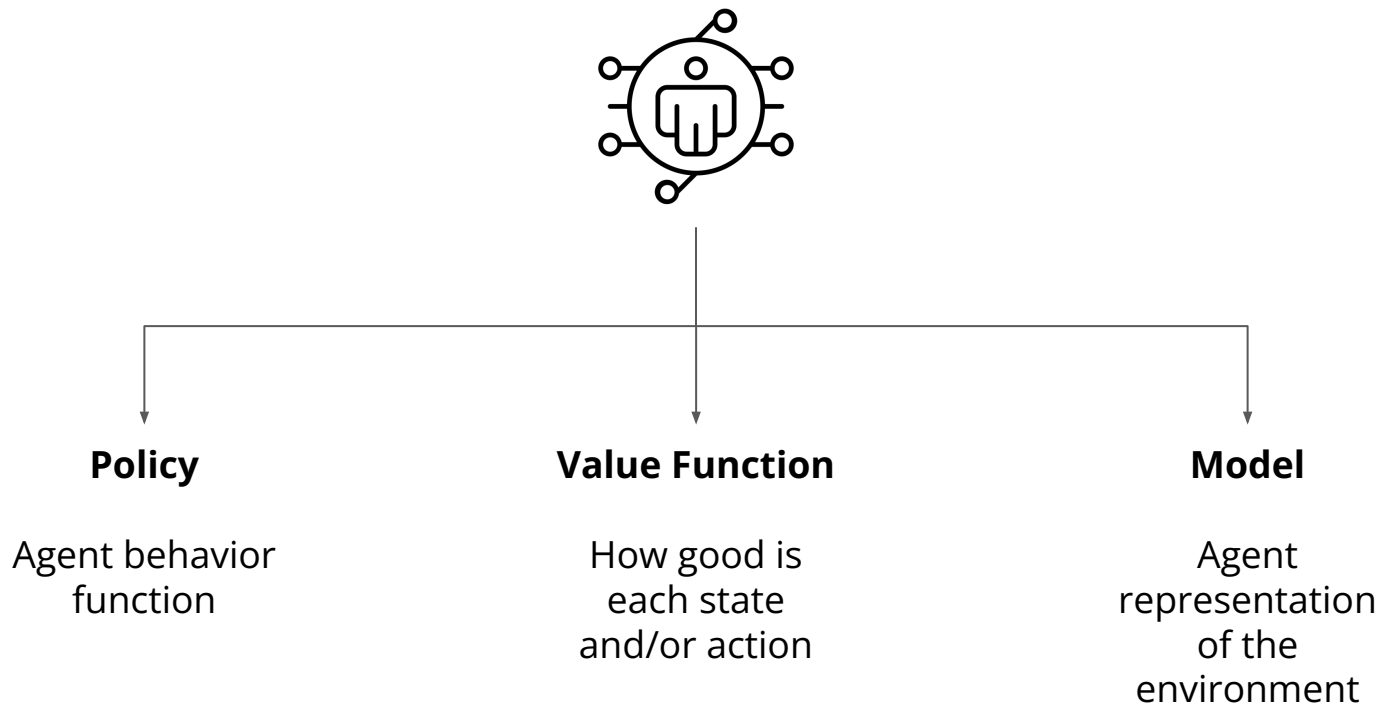
$$O_t = S_t^a = S_t^e$$

- Agent state = environment state = information state
- Formally, this is a Markov Decision Process (MDP)

# PARTIALLY OBSERVABLE ENVIRONMENTS

- **Partial observability:** agent indirectly observes environment:
  - A robot with a camera vision
  - A card player agent
  - A trading agents
- In this case, the agent and the environment state are not the same
- This is a **partially observable Markov decision process (POMDP)**
- An agent must build its own state representation ( $S_a^t$ )

# MAJOR RL AGENT COMPONENTS



# POLICY

- A **policy** is the agent's behaviour
- It is a map from state of action
- There are 2 types of policies:
  - Deterministic policy:  $a = \pi(s)$
  - Stochastic policy:  $\pi(a \mid s) = \mathbb{P}[A_t = a \mid S_t = s]$

# VALUE FUNCTION

- Value function is a prediction of future reward
- We use it to evaluate the goodness of the various states
- We can use it to pick the best action

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$




# MODEL

- A model predicts what the environment will do next
- $P$  predicts the next state
- $R$  predicts the next (immediate) rewards

$$P_{ss'}^a = P[S_{t+1} = s' \mid S_t = s, A_t = a]$$


$$R_s^a = E[R_{t+1} \mid S_t = s, A_t = a]$$

# MAZE EXAMPLE




- Rewards: -1 per time-step
- Action: L,R,U,D
- States: Agent's Location

# MAZE EXAMPLE



- Arrows represent policy  $\pi(s)$  for each state  $s$

# MAZE EXAMPLE



Numbers represent value  $v_{\pi}(s)$  of each state  $s$

# EXPLORATION AND EXPLOITATION

- Reinforcement learning is like trial-and-error learning
- The agents should discover a good policy from its experiences of the environment
- Without losing too much reward along the way
- Exploration find more information about the environment
- Exploitation exploits known information to maximize reward

# RL AGENT TAXONOMY

