

# Markov Decision Process

Mattia Pellegrino, Ph.D Fellow  
[mattia.pellegrino@unipr.it](mailto:mattia.pellegrino@unipr.it)

# Summary

- Markov Processes
- Markov Reward Processes
- Markov Decision Processes

# INTRODUCTION

- Markov decision processes describe formally an environment for reinforcement learning
- Where the environment is fully observable
  - We know everything
- Almost all RL problems can be described as MDP
  - In some ways we can hypothesize all the problems as MDP
  - Partially observable problems can be converted into MDPs

# MARKOV PROPERTY

***An information state (a.k.a. Markov state) contains all useful information from the history***

- ***A state  $S_t$  is a Markov state if and only if***

$$\mathbb{P} [S_{t+1} \mid S_t] = \mathbb{P} [S_{t+1} \mid S_1, \dots, S_t]$$

- The future is independent of the past given present

$$H_{1:t} \longrightarrow S_t \longrightarrow H_{t+1:\infty}$$

- Once the state is known, the history is irrelevant
- The environment state  $S_t^e$  is a Markov State
- The history  $H_t$  is a Markov state

# STATE TRANSITION MATRIX

- For every Markov process we can define the transition probability from a state to another on

$$P_{s,s'} = P[S_{t+1} = s' \mid S_t = s]$$

- State transition matrix  $P$  defines the transition probabilities from all states  $s$  to all successor states  $s'$

$$P = \text{from} \begin{bmatrix} P_{11} & \dots & P_{1n} \\ P_{n1} & \dots & P_{nn} \end{bmatrix}$$

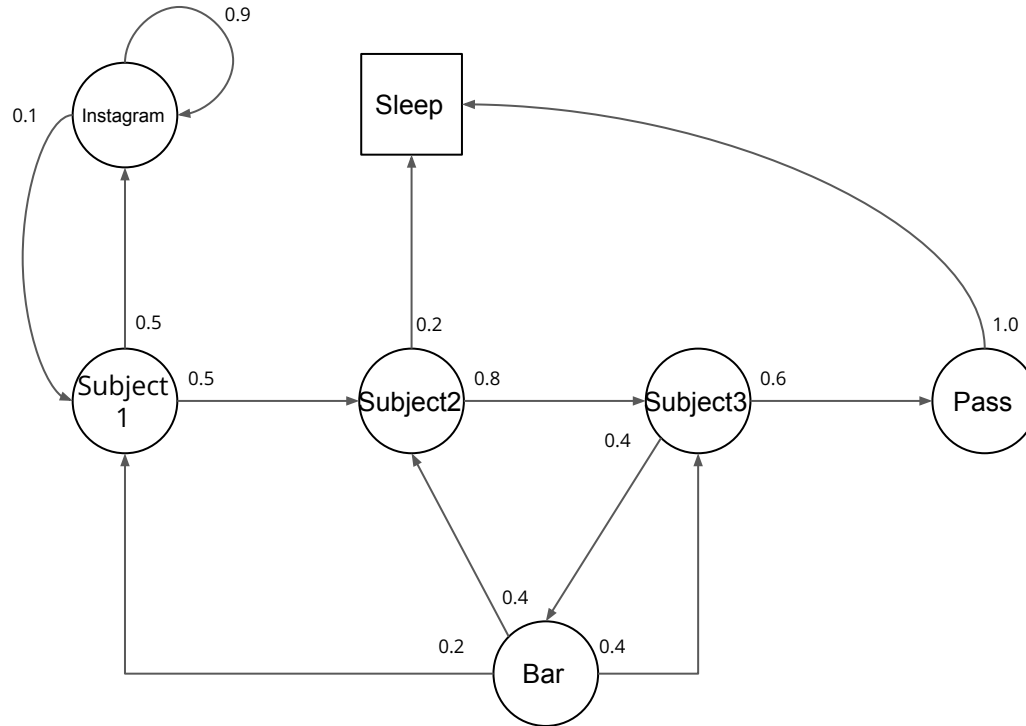
# MARKOV PROPERTY

***A Markov process is a memoryless random process*** (a sequence of random states  $S_1, S_2$  with the Markov property)

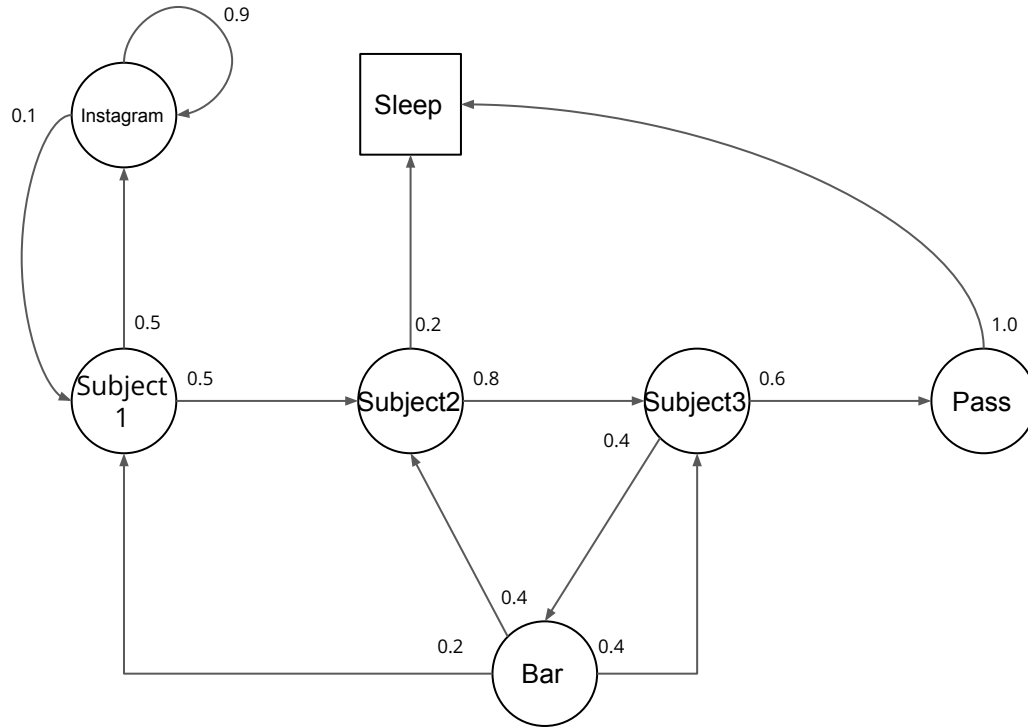
- ***A Markov Process (or Markov Chain) is a tuple  $\langle S, P \rangle$*** 
  - *$S$  is a finite set of states*
  - *$P$  is a state transition probability matrix,*

$$P_{s,s'} = \mathbb{P}[S_{t+1} = s' | S_t = s]$$

# EXAMPLE: STUDENT MARKOV CHAIN




# EXAMPLE: STUDENT MARKOV CHAIN



- Sequence of **episodes** for Student Markov Chain starting from Subject1  
 $S_1, S_2, \dots, S_T$
- S1 S2 S3 Pass Sleep
- S1 IN IN S1 S2 Sleep
- S1 S2 S3 Bar S2 S3 Pass Sleep
- S1 IN IN S1 S2 S3 Bar S1 IN IN  
IN S1 S2 S3 Bar S2 Sleep



# EXAMPLE: STUDENT MARKOV CHAIN



	<i>S1</i>	<i>S2</i>	<i>S3</i>	<i>Pass</i>	<i>Bar</i>	<i>IN</i>	<i>Sleep</i>
<i>S1</i>		0.5				0.5	
<i>S2</i>			0.8				0.2
<i>S3</i>				0.6	0.4		
<i>Pass</i>							1.0
<i>Bar</i>	0.2	0.4	0.4				
<i>IN</i>	0.1					0.9	
<i>Sleep</i>							1

# MARKOV REWARD PROCESS

- A Markov reward process is a Markov chain with values
- **A Markov Reward process is a tuple  $\langle S, P, R, \gamma \rangle$**

- **S is a finite set of states**
- **P is a state transition probability matrix**

$$P_{s',s} = \mathbb{P}[S_{t+1} = s' | S_t = s]$$

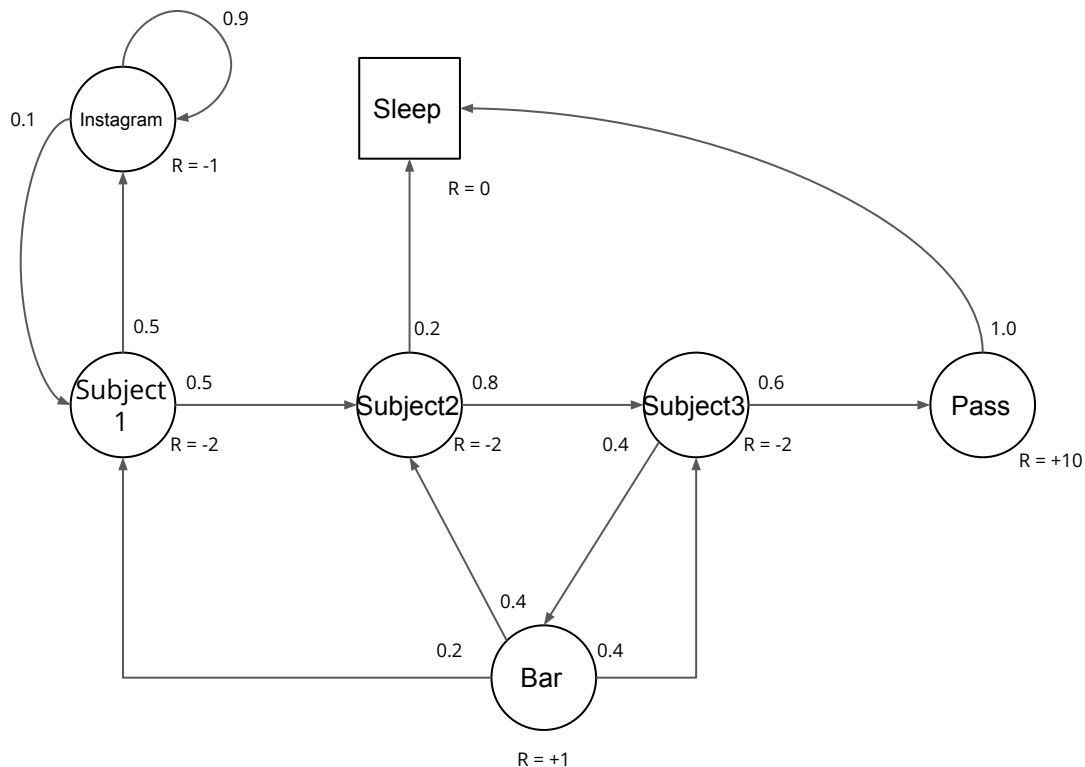
- **R is a reward function**

$$R_s = \mathbb{E}[R_{t+1} | S_t = s]$$

- **Gamma is a discount factor**

$$\gamma \in [0, 1]$$

# EXAMPLE: STUDENT MARKOV CHAIN



# RETURN

- The return  $G_t$  (Goal) is the total discounted reward from time-step  $t$

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- The discount  $\gamma \in [0, 1]$  is the present value of the future rewards
- The value of receiving  $R$  after  $k+1$  timeteps is  $\gamma^{k+1} R$
- if  $\gamma$  is close to 0 than we have a myopic evaluation, if is close to 1 is far-sighted evaluation

# VALUE FUNCTION

- The value function  $v(s)$  of an MRP is the expected return starting from state  $s$

$$v(s) = \mathbb{E}[G_t | S_t = s]$$

# EXAMPLE: STUDENT MRP

- Starting from  $S_1 = S1$  with  $\gamma = 1/2$

**S1 S2 S3 Pass Sleep**

$$v1 = -2 - 2 * 1/2 - 2 * 1/4 + 10 * 1/8 = -2.25$$

**S1 IN IN S1 S2 Sleep**

$$v1 = -2 - 1 * 1/2 - 1 * 1/4 - 2 * 1/8 - 2 * 1/16 = -3.125$$

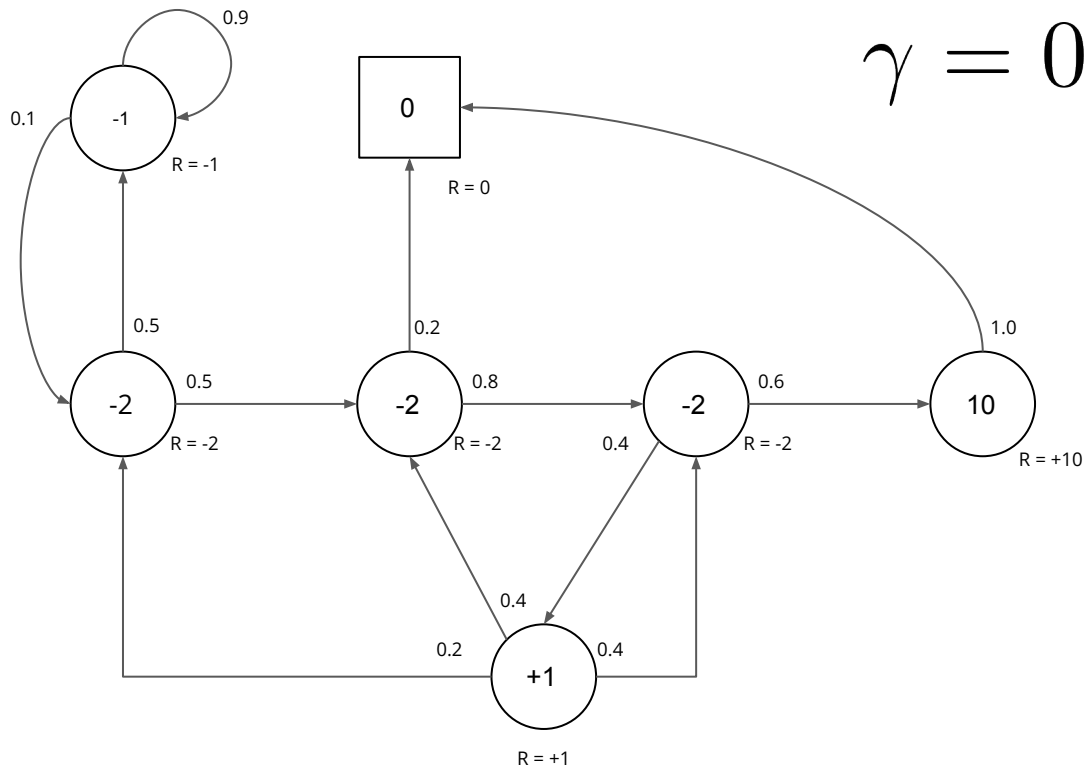
**S1 S2 S3 Bar S2 S3 Pass Sleep**

$$v1 = -2 - 2 * 1/2 - 2 * 1/4 + 1 * 1/8 - 2 * 1/16 \dots = -3.41$$

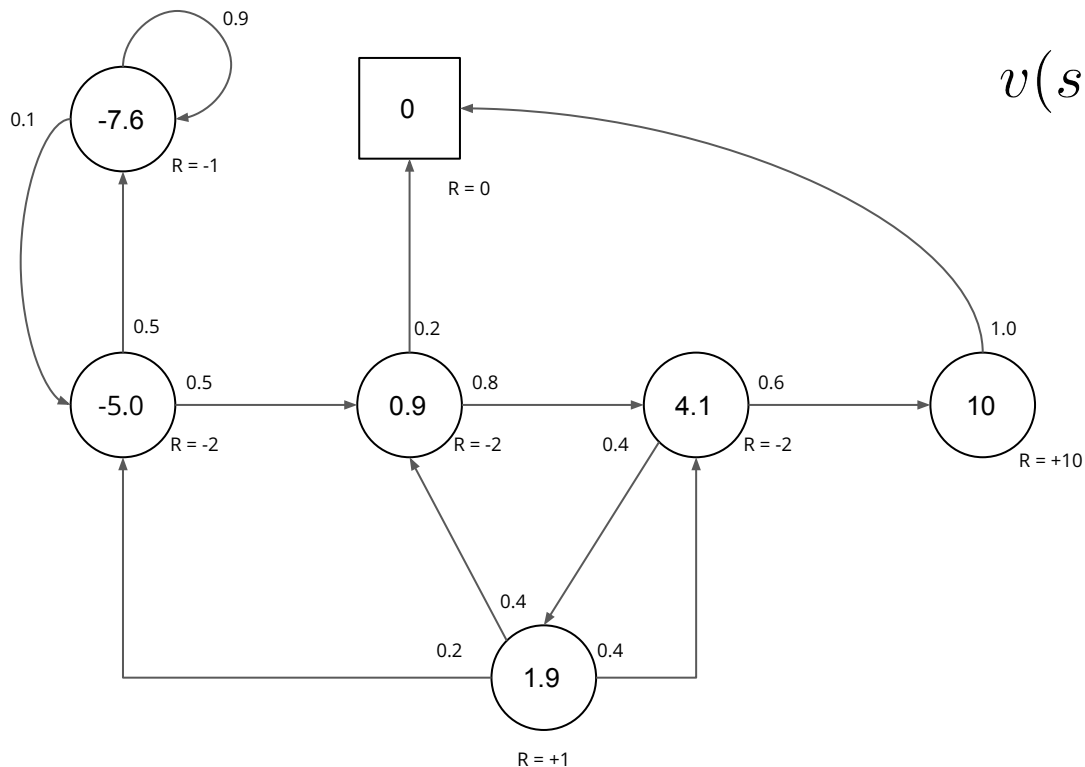
**S1 IN IN S1 S2 S3 Bar**

$$v1 = -2 - 1 * 1/2 - 1 * 1/4 - 2 * 1/8 - 2 * 1/16 \dots = -3.20$$

# EXAMPLE: STATE-VALUE FUNCTION (1)

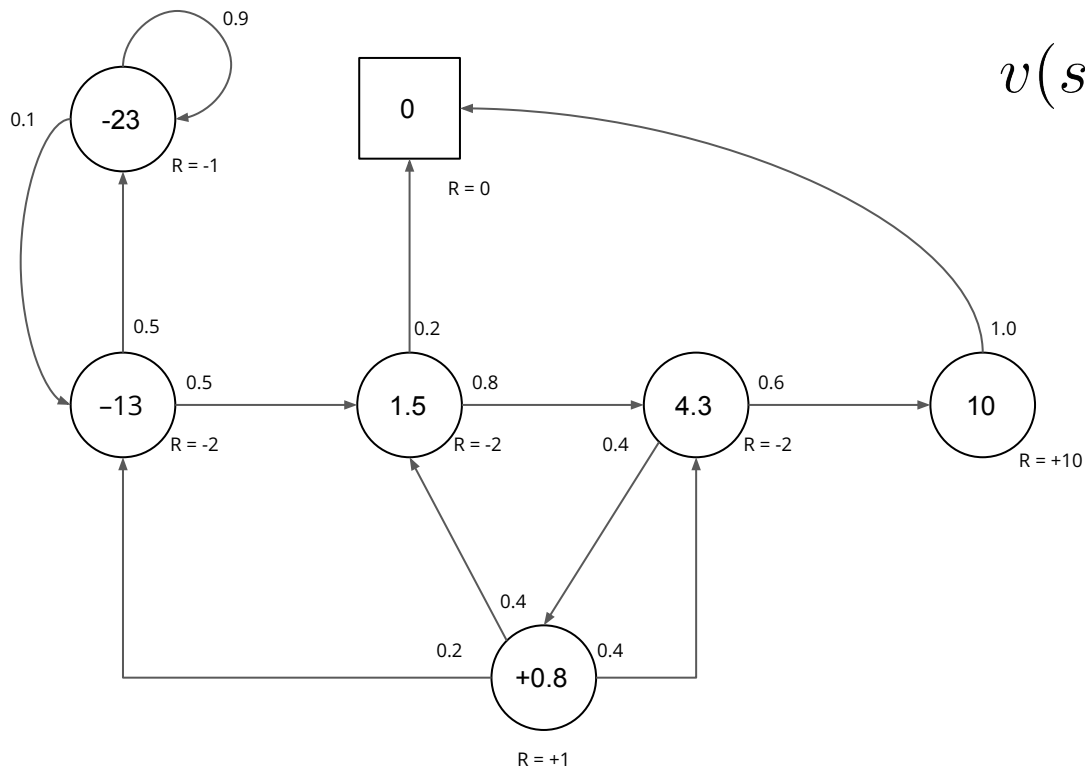


# EXAMPLE: STATE-VALUE FUNCTION (2)





# EXAMPLE: STATE-VALUE FUNCTION (3)




# BELLMAN EQUATION

- The value function can be divided into two parts:
  - Immediate reward  $R_{t+1}$
  - Discounted value of successor state

$$\begin{aligned} v(s) &= \mathbb{E} [G_t \mid S_t = s] \\ &= \mathbb{E} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E} [R_{t+1} + \gamma (R_{t+2} + \gamma R_{t+3} + \dots) \mid S_t = s] \\ &= \mathbb{E} [R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \mathbb{E} [R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s] \end{aligned}$$

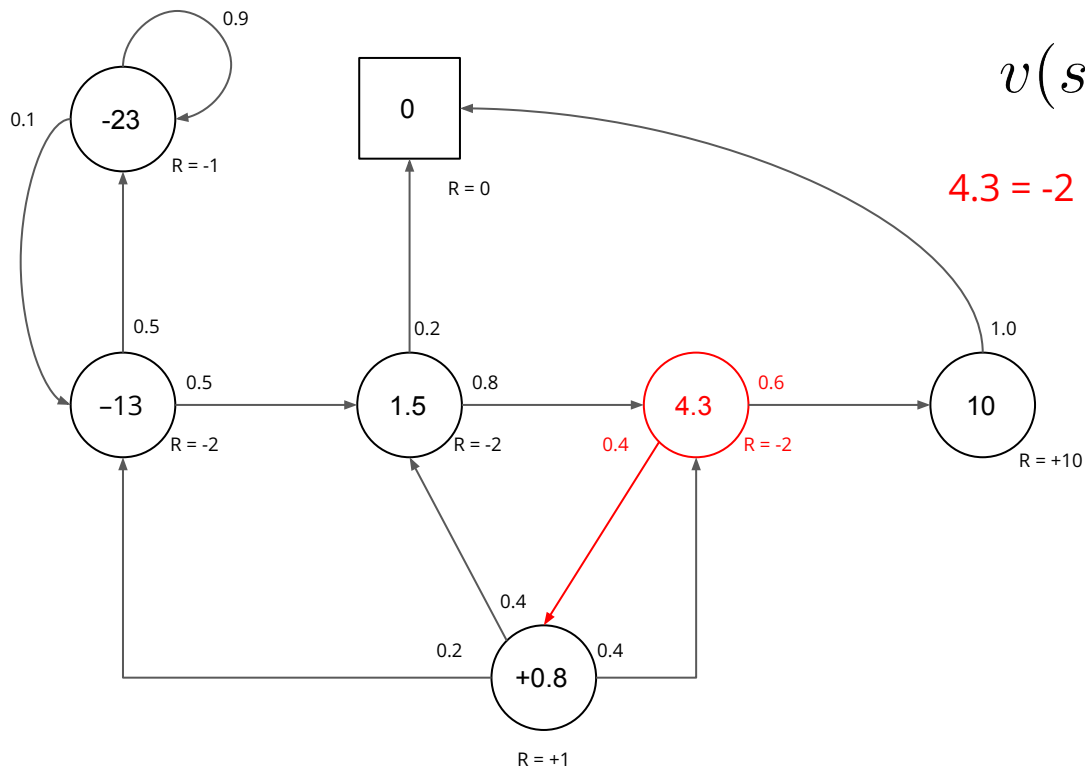
# BELLMAN EQUATION

$$v(s) = \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$$



$$v(s) = R_s + \gamma \sum_{s' \in S} P_{ss'} v(s')$$

# EXAMPLE: BELLMAN EQUATION



$$v(s), \gamma = 1$$

$$4.3 = -2 + 0.6 * 10 + 0.4 * 0.8$$

# BELLMAN EQUATION (MATRIX)

- The Bellman equation can be converted into a matrix form

$$v = R + \gamma P v$$

Where  $v$  is a column vector in which every row is a state

$$\begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix} = \begin{bmatrix} \mathcal{R}_1 \\ \vdots \\ \mathcal{R}_n \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \vdots & & \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix} \begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix}$$

# SOLVING THE BELLMAN EQUATION

- The Bellman equation is linear and can be solved

$$v = R + \gamma P v$$

$$(I - \gamma P)v = R$$

$$v = (I - \gamma P)^{-1} R$$

- Direct solution only possible for small MRPs
- There are many iterative methods for large MDPs:
  - Dynamic programming
  - Monte-Carlo evaluation
  - Temporal-Difference learning

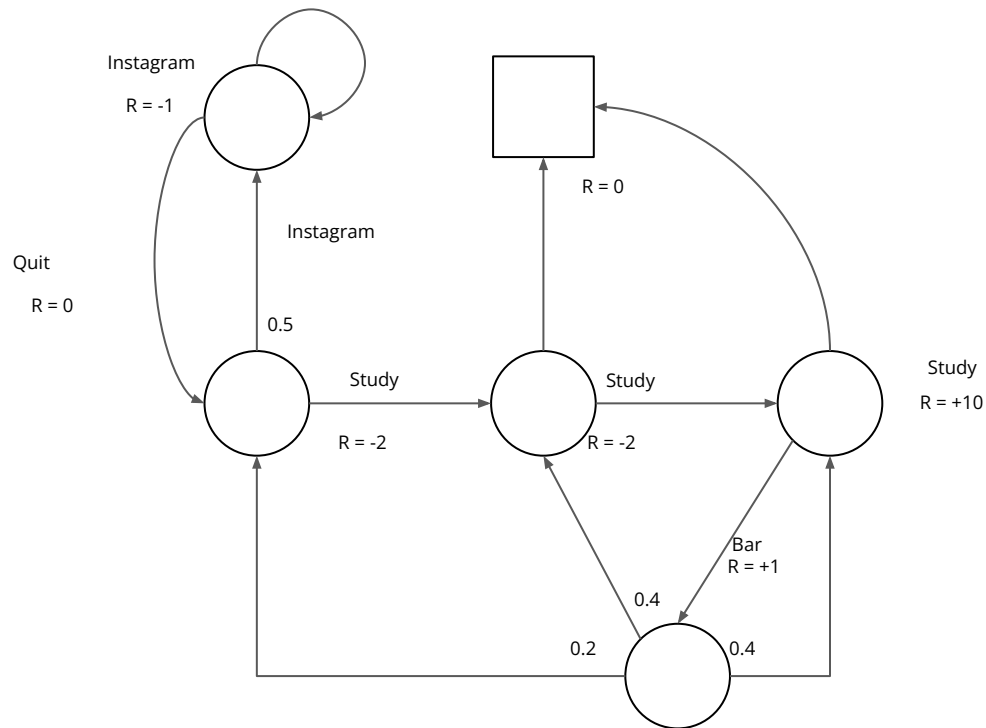
# MARKOV DECISION PROCESS

- **A Markov Decision process (MDP) is a Markov reward process with decision. All states are Markov in the environment.**

A Markov decision Process is a tuple  $\langle S, A, P, R, \gamma \rangle$

- $S$  is a finite set of states
- $A$  is a finite set of actions
- $P$  is a state transition probability matrix,  
 $P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- $R$  reward function,  $R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
- $\gamma$  is a discount factor  $\gamma \in [0, 1]$

# EXAMPLE: STUDENT MDP





# POLICIES

- A policy  $\pi$  is a distribution over actions given states,

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$$

- A policy defines (in a fully way) the agent's behaviour
- MDP policies depend on the current state (we throw away the history)
- The policies are stationary and time-independent

# POLICIES

- Given an MDP  $M = \langle S, A, P, R, \gamma \rangle$  and a policy  $\pi$ :
  - The state sequence  $S_1, S_2 \dots$  is a Markov process  $\langle S, P^\pi \rangle$
  - The state and the reward sequence  $S_1, R_1, S_2 \dots$  is a Markov reward process  $\langle S, P^\pi, R^\pi, \gamma \rangle$

Where:

$$P_{s,s'}^\pi = \sum_{a \in A} \pi(a|s) P_{s,s'}^a$$

$$R_s^\pi = \sum_{a \in A} \pi(a|s) R_s^a$$

# VALUE FUNCTION

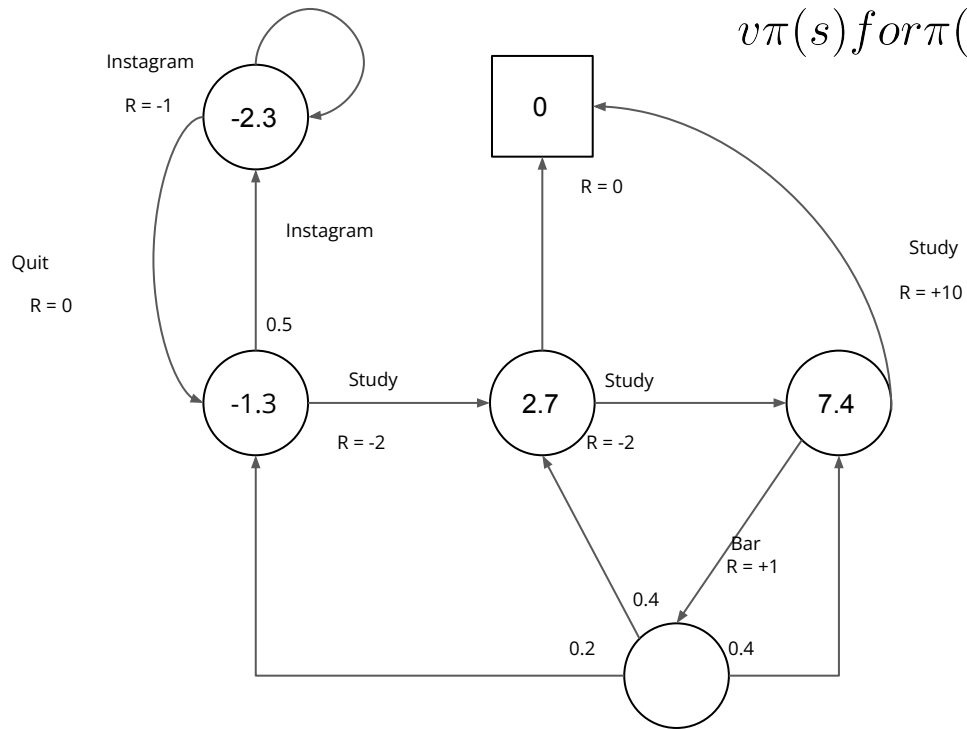
- The state-value function  $v_{\pi}(s)$  of an MDP is the expected return starting from state  $s$ , and then following the policy

$$v_{\pi} = \mathbb{E}_{\pi}[G_t | S_t = s]$$

- The action-value function  $q_{\pi}(s, a)$  is the expected return starting from state  $s$ , taking action  $a$ , according to policy  $\pi$

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

## EXAMPLE: STUDENT MDP



$$v\pi(s)for\pi(a|s) = 0.5, \gamma = 1$$

# BELLMAN EXPECTATION EQUATION

- The state-value function can be split into immediate reward plus the discounted value of successor state


$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]$$

- The action-value function can similarly be splitted

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$


# BELLMAN EXPECTATION EQUATION

$$V_{\pi}$$




$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a)$$

$$Q_{\pi}$$




$$q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s')$$

# BELLMAN EXPECTATION EQUATION - 2

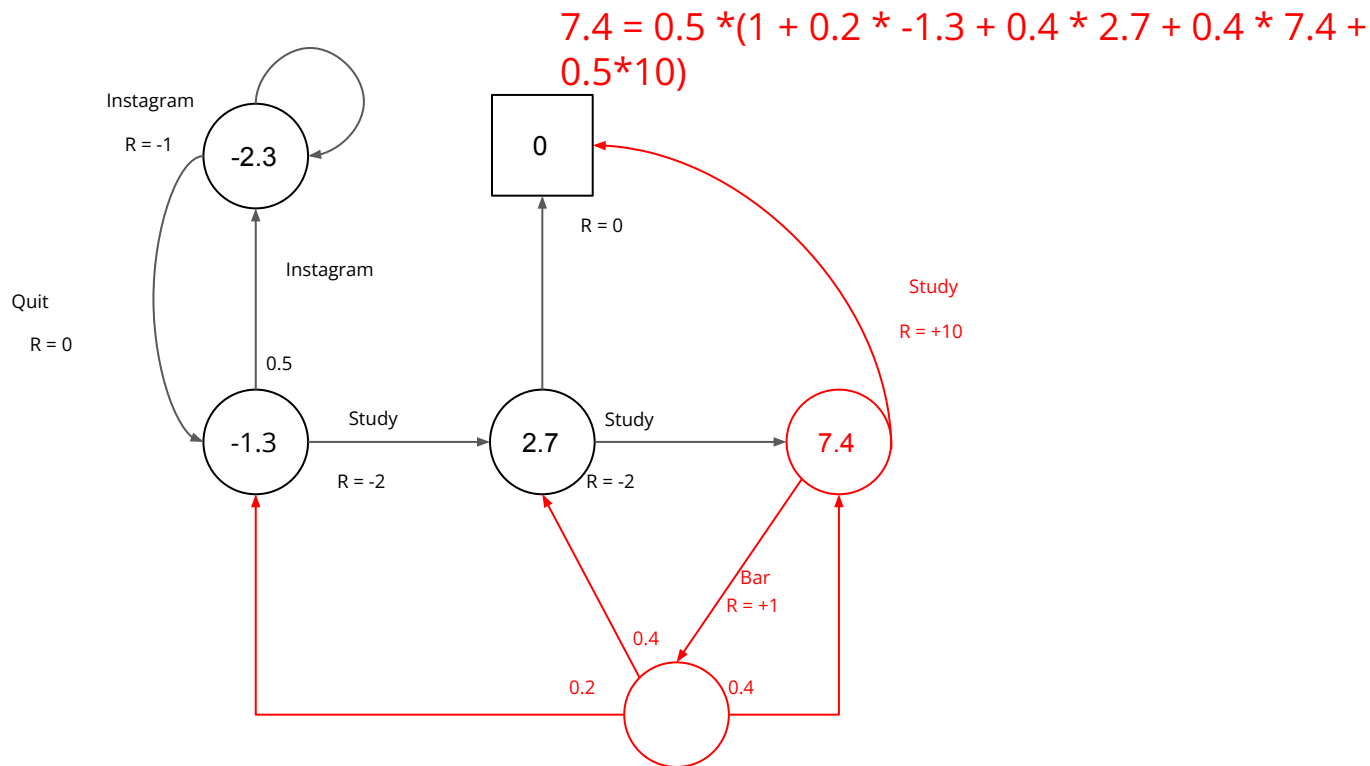


$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a | s) (\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s'))$$



$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a' | s') q_\pi(s', a')$$

# EXAMPLE: STUDENT MDP





# OPTIMAL VALUE FUNCTION

- The optimal state-value function  $v_*(s)$  is the maximum value function over all policies

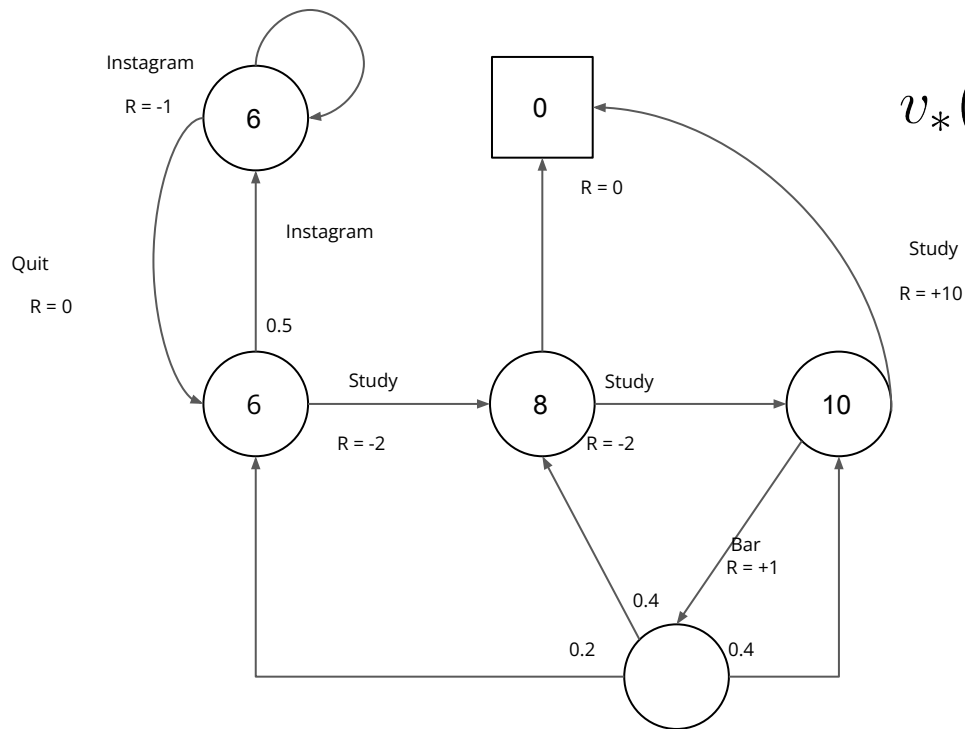
$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- The optimal action-value function  $q_*(s,a)$  is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

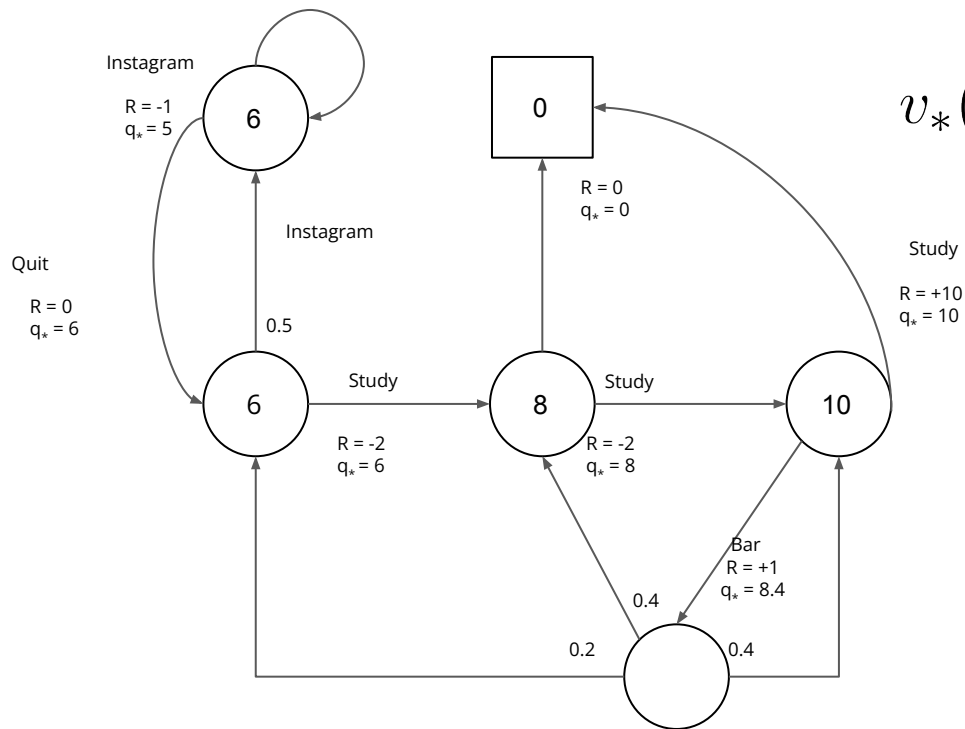
- This function specifies the best possible performance in the MDP.
- An MDP is “solved” when we discover the optimal value

# EXAMPLE: STUDENT MDP (optimal value function)



$v_*(s)$  for  $\gamma = 1$

# EXAMPLE: STUDENT MDP (optimal action-value function)



$v_*(s)$  for  $\gamma = 1$

# OPTIMAL POLICY

- For any Markov Decision Process
  - There is an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi, \forall \pi$
  - All optimal policies achieve the optimal value function,  $v_{\pi_*} = v_*(s)$
  - All optimal policies achieve the optimal action-value function,  $q_{\pi_*}(s, a) = q_*(s, a)$

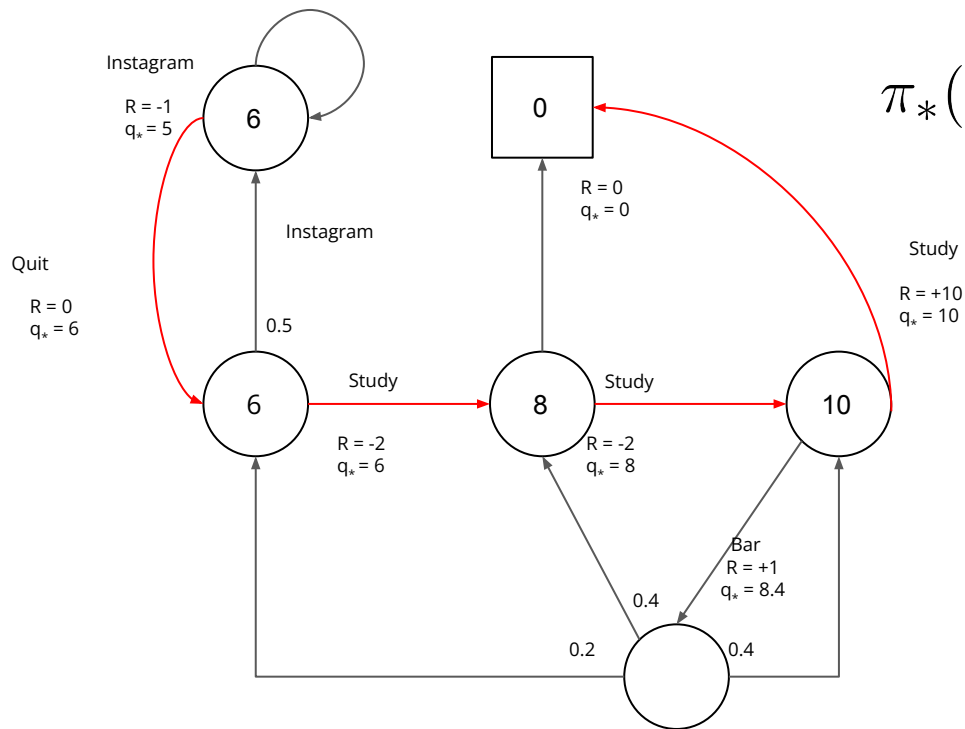
# OPTIMAL POLICY

- An optimal policy can be found maximizing over  $q_*(s, a)$

$$\pi_*(a \mid s) = \begin{cases} 1 & \text{if } a = \underset{a \in \mathcal{A}}{\operatorname{argmax}} q_*(s, a) \\ 0 & \text{otherwise} \end{cases}$$

- There is always an optimal policy (deterministic) for any MDP
- If we know  $q_*(s, a)$ , we already have the optimal policy


# EXAMPLE: STUDENT MDP (optimal policy)



$$\pi_*(a|s) \text{ for } \gamma = 1$$


# BELLMAN OPTIMALITY EQUATION

$V_*$




$$v_* = \max_a q_*(s, a)$$

$Q_*$




$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s')$$

# BELLMAN OPTIMALITY EQUATION - 2




$$v_*(s) = \max_a R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s')$$



$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \max_{a'} q_*(s', a')$$



# EXAMPLE: STUDENT MDP (optimal policy)

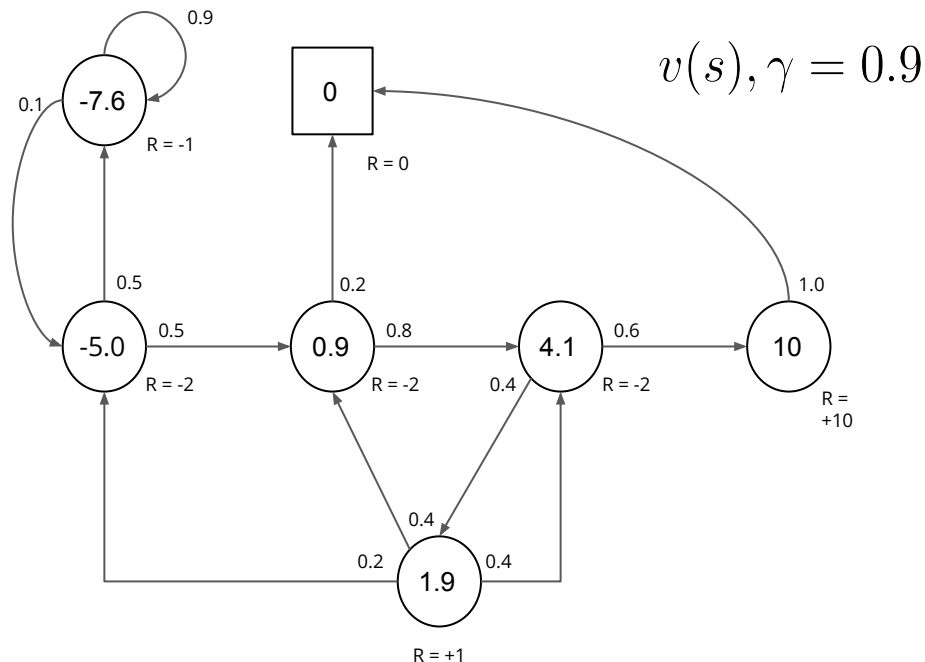


# SOLVING BELLMAN OPTIMALITY EQ

- The Bellman Optimality Equation
  - Is not linear
  - No closed form solution
  - Can be solved with many iterative solution methods:
    - Value Iteration
    - Policy Iteration
    - Q-learning
    - Sarsa

# ASSIGNMENT: BUILDING A STUDENT DECISION-MAKING PROBLEM

- **Objective:** build a graph representing a **student decision-making problem**, compute **value functions** using the **Bellman equation**, and **visualize** the **graph with annotated value functions**.



# ASSIGNMENT: BUILDING A STUDENT DECISION-MAKING PROBLEM

- **1 - Building the graph:**
  - Import necessary libraries: networkx, matplotlib.pyplot, numpy, seaborn, and pyvis.network.
  - Define the states and rewards for each state.
  - Define the transition probabilities between states.
  - Create a directed graph using nx.DiGraph().
  - Add nodes to the graph for each state.
  - Add edges to represent transitions between states, including associated actions and rewards.

# ASSIGNMENT: BUILDING A STUDENT DECISION-MAKING PROBLEM

- **2 - Visualizing Transition Probabilities:**
  - Compute the transition matrix based on the defined transition probabilities.
  - Plot the transition matrix using `seaborn.heatmap()`.
- **3 - Computing Value Functions:**
  - Define the discount factor.
  - Use the Bellman equation to compute the value function of each state.

# ASSIGNMENT: BUILDING A STUDENT DECISION-MAKING PROBLEM

- **4 - Visualizing the Graph with Value Functions:**
  - Draw the graph using `nx.draw()` with appropriate node positions.
  - Annotate each node with its corresponding value function.
  - Annotate each edge with associated actions and rewards.
- **5 - Optional: Visualizing Value Function Dynamics:**
  - Use `matplotlib` to visualize how the value function changes with different discount factors.

# ASSIGNMENT: BUILDING A STUDENT DECISION-MAKING PROBLEM

- **Submission:**

- Submit the Python script containing the code to construct the graph, compute value functions, and visualize the graph with annotated value functions.
- <https://shorturl.at/gjyO7> - Expiration date: 25th May 2024

- **Note:** Ensure the code is well-commented and understandable. Avoid using external guidance or assistance, as the purpose of the assignment is to demonstrate understanding and implementation independently.